

CENG 322
LAB 9


Kernel Code

Kernel

- A kernel is a **central component** of an operating system.
- It acts as an **interface** between the user applications and the hardware.
- The main aim of the kernel is to **manage the communication** between the **software** (user level applications) and the **hardware** (CPU, disk memory etc).

The main tasks:

- Process management
- Device management
- Memory management



Kernel Types

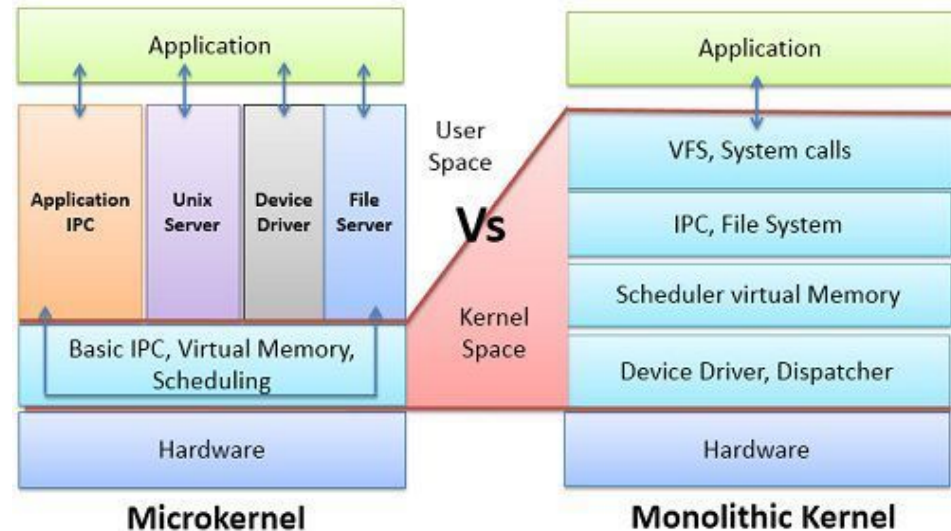
Anciently, all the basic **system services** like **process** and **memory** management, **interrupt** handling etc. were packaged **into a single module** in kernel space.

This type of architecture led to some serious **drawbacks**:

- Size of kernel, which is **huge**.
- Poor maintainability (i.e. bug fixing or addition of new features resulted in **recompilation of the whole kernel** which could consume hours)

What are the types? (Mainly)

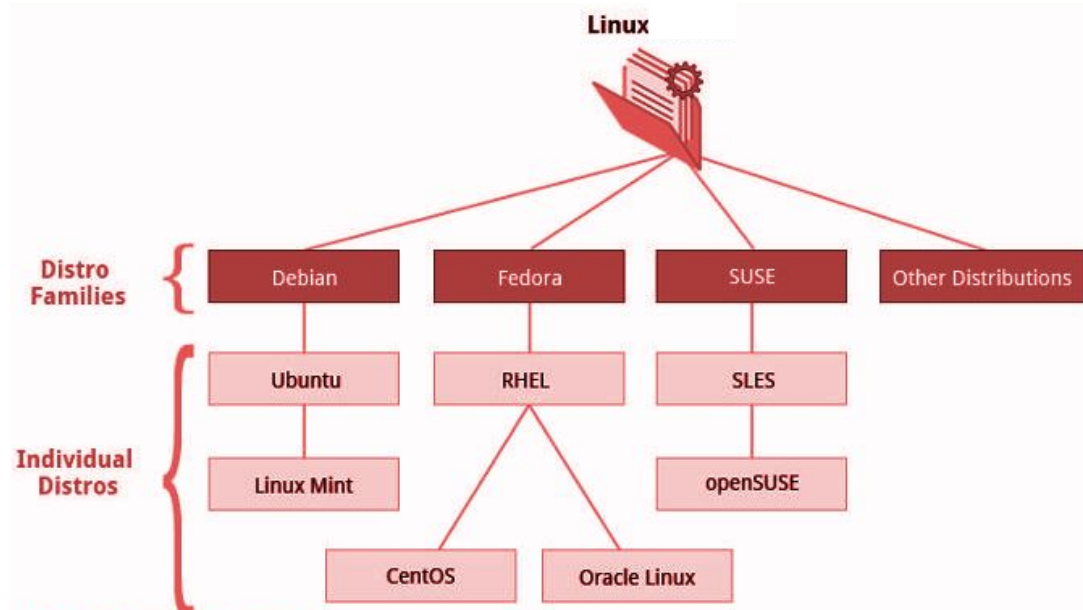
- **Monolithic Kernel**: OS runs as a single program in KM. (Linux)
- **Micro Kernel**: User and kernel services are implemented in different address space. (Mac OS X)



Linux

Is Linux a kernel?

What is Linux distribution?



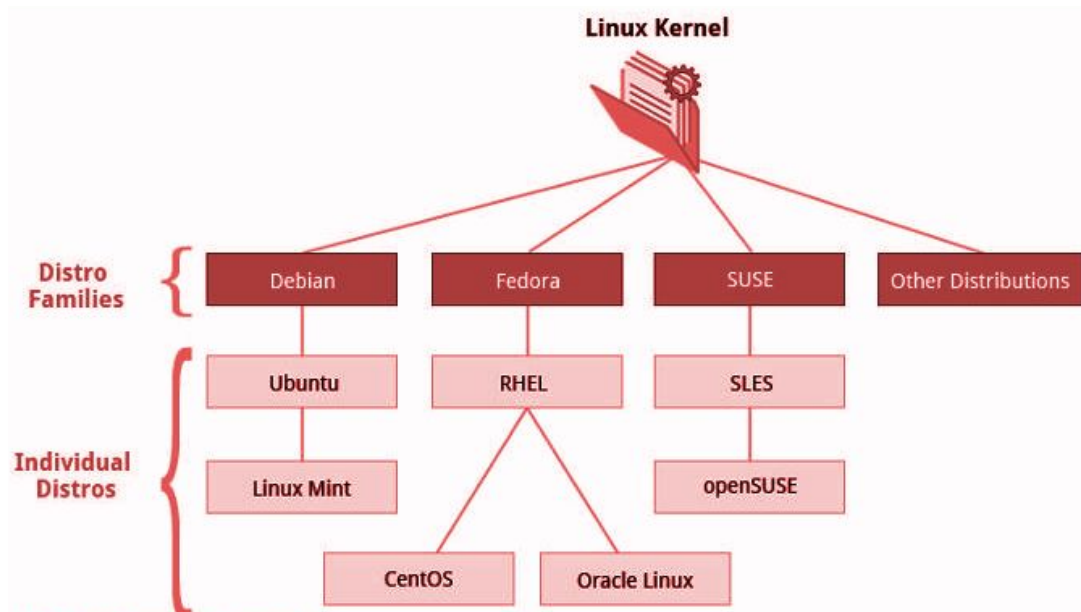
Linux

Is Linux a kernel?

- **YES**. Linux is just a kernel, not an OS!

What is Linux distribution?


- It is an OS composed of the Linux **kernel**, GNU **tools**, additional **software** and a package **manager**.
- Kernel + Additional Programs



Linux

Why are Linux distributions **UNIX-like** OS?

- Before Linux kernel, Richard Stallman created the GNU project, the first free software project, in 1983.
- The GNU project implemented many of the popular Unix utilities like cat, grep, ls, shell (bash) along with developing their own compilers (GCC) and editors (Emacs).
- Back in the 80s UNIX was super expensive, thereof Linus Torvalds developed a new kernel that was like UNIX. With the GNU tools (especially shell commands), it is behaved like UNIX.
- Check source code: <https://github.com/torvalds/linu>




Languages







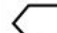
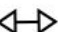
Kernel Module


- Modules are **pieces of code** that can be **loaded** and **unloaded** into the kernel.
- They extend the **functionality** of the kernel without the need to **reboot the system**. (e.g. **device driver** allows the kernel to access hardware connected to the system)
- Without modules, we would have to add new functionality **directly into the kernel image** (requires **rebuild** and **reboot** the kernel).

User Space:



Legend:

	— Process		— User-Kernel Communications (system calls, signals)
	— Code Module		— Inter-Kernel Communications (kernel API)
	— Hardware		— Hardware Communications (interrupts, ports I/O)



Writing Kernel Module

Kernel Code

- Risk data loss and **system corruption** in the case of a any kernel code mistake !
- If you have a fault, it will **lock up the entire system**.
- Using a **virtual machine** is safer and eliminates risks.
- No access to the **standard library** (e.g. printf - replacement for printk and kmalloc - similar to malloc).
- There is **no garbage collection**. After unloading module, we have to **cleaning up**.
- Required **root privileges**.

Kernel Module

- **Kernel modules** typically act as **APIs** for user space programs.
- To discover **what modules** are already loaded within your current kernel use the command **lsmod**.
- For information about module: **modinfo**.
- Loaded modules are stored within the file **/proc/modules**, check: `sudo cat /proc/modules`
- To **search** for the video module: `sudo lsmod | grep video`
- **Built-in** kernel modules: `more /lib/modules/$(uname -r)/modules.builtin | head -10`




Writing Kernel Module

- **Static Modules**

- Those are compiled as **part of the base kernel** and it is **available** at any time (compiled in).
- If I compile “asus” **device driver** as a static module along with the kernel, I will be adding an **extra size to the kernel image** permanently. Suppose if I **modify** “asus” device driver then need to **re-compile** the entire kernel to build it. Every time needs to **rebooted** for the changes to take effect.

- **Dynamic Modules**

- Those are compiled as modules **separately** and **loaded** based on user demand. These are also called as **Loadable Kernel Modules(LKM)**.
- It relies on the presence of libraries, and the library is **only loaded once** no matter how many programs use it.
- If we build “asus” module as dynamically(LKM), do **not need to rebuild** the kernel. It can be compiled separately. It can be loaded into the kernel at runtime without having the machine to reboot.



Writing Kernel Module

Dynamic Module Commands

- **insmod**: Inserting a module into the kernel. Kernel object (suffixed with .ko) can be installed with an option.

```
$ insmod my_module.ko my_option=1
```

- **rmmod**: Removing a module. The kernel will not remove a module currently being used, -w option is removing once the use-count has decreased to zero.


```
$ rmmod my_module
```

- **modprobe**: Intelligent version of insmod and rmmod, it considers dependencies (if that particular module is dependent on any other module) and loads them.

```
$ modprobe module_name parameter=value
```

- **dmesg**: displays kernel-related messages.

```
$ dmesg
```



Writing Kernel Module

- **Step 1/4:** Checking kernel version and installing required packages.

- Check kernel version:

```
sudo uname -a
```

- Install the essential development tools and the kernel headers necessary for module development:

```
apt-get install build-essential linux-headers-$(uname -r)
```

- User space header(s) can not be included in this module (e.g. `stdio.h`, `stdlib.h...`) headers.
- Linux kernel headers are located under `usr/include/linux/`.

Writing Kernel Module

- **Step 2/4:** Creating a code file (.c).

```
#include <linux/init.h> // for the macros
#include <linux/module.h> // used by all modules
#include <linux/kernel.h> // for KERN_INFO
```

```
MODULE_LICENSE("GPL");
MODULE_AUTHOR("A. Yigit");
MODULE_DESCRIPTION("A simple example Linux module.");
MODULE_VERSION("0.01");
```

Check: CENG322_module.c

Writing Kernel Module

- **Step 2/4:** Creating a code file (.c).
- Printk displays the message in the kernel log (/var/log/kern.log).

```
#define KERN_EMERG    "<0>" /* system is unusable */
#define KERN_ALERT    "<1>" /* action must be taken immediately */
#define KERN_CRIT     "<2>" /* critical conditions */
#define KERN_ERR      "<3>" /* error conditions */
#define KERN_WARNING  "<4>" /* warning conditions */
#define KERN_NOTICE   "<5>" /* normal but significant condition */
#define KERN_INFO     "<6>" /* informational */
#define KERN_DEBUG    "<7>" /* debug-level messages */
```

```
static int  init my_init(void) {
    printk(KERN_INFO "Hello, World! \n");
    return 0;
}

static void  exit my_exit(void) {
    printk(KERN_INFO "Goodbye, World! \n");
}

module_init(my_init);
module_exit(my_exit);
```

Check: CENG322_module.c

Writing Kernel Module

- **Step 3/4:** Compiling the code using a makefile.

```
obj-m += CENG322_module.o
all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

- **obj-y:** static compilation will link to kernel image.
obj-m: dynamic compilation can be loaded into kernel image.
- Kernel object(.ko) will be created after compiling.

Check: Makefile

Writing Kernel Module

- **Step 4/4:** Inserting the module into the kernel.

```
sudo insmod CENG322_module.ko
```

- printk function doesn't output to the console but rather the kernel log

```
sudo dmesg
```

- Checking the kernel modules again:

```
lsmod | grep "CENG322"
```

- Removing the kernel:

```
sudo rmmod CENG322_module
```

- If you run dmesg again, you'll see "Goodbye, World!" in the logs. You can also use lsmod again to confirm it was unloaded.

Passing Arguments

- Declare the variables as `global` and then use the `module_param()` macro, (defined in `include/linux/moduleparam.h`).
- At runtime, `insmod` will fill the variables, like `insmod mymodule.ko myvariable=5`.
- The `module_param()` macro takes **3 arguments**; name, type and permissions for the corresponding file in `sysfs`.
- For arrays of integers or strings see `module_param_array()` and `module_param_string()`.

```
int myint = 3;
module_param(myint, int, 0);

int myintarray[2];
module_param_array(myintarray, int, NULL, 0);

short myshortarray[4];
int count;
module_param_array(myshortarray, short, &count, 0);
```

- **Check:** `CENG322_modparameter.c`

Passing Arguments

- You can find modes in the `<linux/stat.h>` header.
 - `S_IRUSR, S_IWUSR, S_IXUSR` – owner permission
 - `S_IRGRP, S_IWGRP, S_IXGRP` – group permission
 - `S_IROTH, S_IWOTH, S_IXOTH` – other permission
- Check all permissions:
https://www.gnu.org/software/libc/manual/html_node/Permission-Bits.html

Communication with User App

- Write a kernel module that shows a message if user space application **reads data from** /proc.
- Message will be "You read data!".
- Overriding read function.

```
static struct proc_dir_entry* proc_entry;

static ssize_t custom_read(struct file* file, char __user* user_buffer,
size_t count, loff_t* offset)
{
    printk(KERN INFO "Calling our own custom read method.");
    char message[] = "You read data!\n";
    int message_length = strlen(message);
    if (*offset > 0)
        return 0;
    copy_to_user(user_buffer, message, message_length);
    *offset = message_length;

    return message_length;
}
```

- **Check:** CENG322_comuserapp.c

Communication with User App

```
static struct file_operations fops =
{
    .owner = THIS_MODULE,
    .read = custom_read
};

// Custom init and exit methods
static int __init custom_init(void) {
    proc entry = proc_create("iytecommodule", 0666, NULL, &fops);
    printk(KERN_INFO "Communication driver loaded.");
    return 0;
}
```

```
static void __exit custom_exit(void) {
    proc remove(proc entry);
    printk(KERN_INFO "Goodbye ...");
}

module_init(custom_init);
module_exit(custom_exit);
```

Developing a Device Driver

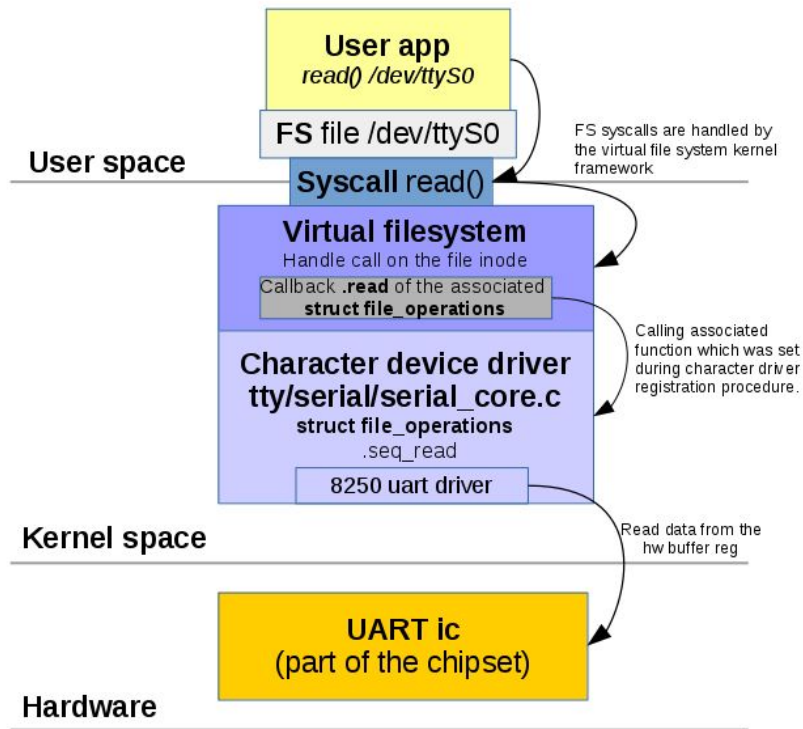
- **Drivers** are used to help the hardware devices interact with OS.
- Linux works with an “**everything is a file**” abstraction.
- Hardware **devices** are treated like **ordinary files**, which makes it **easier** for the software to **interact** with the device drivers.
- When a device is connected to the system, a device file is created in **/dev** directory.
- Most Common types of devices in Linux:
 - **Character devices**: These devices transmit the data character by characters, like a mouse or a keyboard (no buffering, sequentially).
 - **Block devices**: These devices transfer unit of data storage called a block, USB drives, hard drives, and CD ROMs (randomly).



Developing a Device Driver

```
ls -l /dev/ttyS0
```

- Symbol **C**, in the beginning, means that this device is a **character device**.
- **Major** and **Minor** numbers are assigned to the device.
- Inside the Linux kernel, every device is identified **not by a symbolic name** but by **a unique number** – the device's major number.
- This number assigning by the kernel during device **registration**. Every device driver can support multiple “sub-devices”.
- A serial port adapter may contain **two hardware ports**. Both of these ports are handled by the same driver, and they **share one Major number**. But inside this driver, **each of these ports** is also identified by the unique number, and this is a device **Minor** number.



Developing a Device Driver

- Create a **read-only char device** that says **how many times you have read** from the dev file. **Writing** will not be supported !
- Simply read in the data and **print a message** acknowledging that we received it.
- In the **multiple-threaded** environment, without any protection, concurrent access to the same memory may lead to the **race condition**.
- In the kernel module, this problem may happen due to **multiple instances** accessing the shared resources.
- We can use atomic Compare-And-Swap (CAS) to maintain the states, **CDEV_NOT_USED** and **CDEV_EXCLUSIVE_OPEN**, to determine whether the file is **currently opened** by someone or not.
-
- CAS compares the **contents of a memory location** with the expected value and, only if they are the same, modifies the contents of that memory location to the desired value.

```
enum {  
    CDEV_NOT_USED = 0,  
    CDEV_EXCLUSIVE_OPEN = 1,  
};  
  
/* Is device open? Used to prevent multiple access to device */  
static atomic_t already_open = ATOMIC_INIT(CDEV_NOT_USED);
```

- **Check:** CENG322_devdriver.c

Kernel Building and Custom System Call

- All Linux distributions are based on a predefined kernel.
- In order to **disable** several options, drivers and **change** kernel **settings**, we need to build **kernel** from scratch.
- It takes a lot of time (hours).
- Firstly download source code of Linux kernel.

```
wget https://cdn.kernel.org/pub/linux/kernel/v5.x/linux-5.17.4.tar.xz
```


- **Check:** custom_system_call.txt

The Linux Kernel Archives

[About](#) [Contact us](#) [FAQ](#) [Releases](#) [Signatures](#) [Site news](#)

Protocol
[HTTP](#)
[GIT](#)
[RSYNC](#)

Location
<https://www.kernel.org/pub/>
<https://git.kernel.org/>
[rsync://rsync.kernel.org/pub/](https://rsync.kernel.org/pub/)

Latest Release
5.17.4 

mainline:	5.18-rc3	2022-04-17	[tarball]	[patch]	[inc. patch]	[view diff]	[browse]	
stable:	5.17.4	2022-04-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]	
stable:	5.16.20 [EOL]	2022-04-13	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]	
longterm:	5.15.35	2022-04-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]	
longterm:	5.10.112	2022-04-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]	
longterm:	5.4.190	2022-04-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]	
longterm:	4.19.239	2022-04-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]	
longterm:	4.14.276	2022-04-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]	
longterm:	4.9.311	2022-04-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]	
linux-next:	next-20220422	2022-04-22						[browse]

Thanks ...